



Discussion:

The Mean Threshold Graph does not clearly show it, but the slow and fast estimations of p^* are almost exact for each N value. This is as expected since the p^* value should not be based on the speed at which the program recognizes percolation, but on the number of opened slots.

In contrast to the first graph, the Mean Running Time per Experiment Graph shows a clear difference between the fast and slow implementations. As is fancily named, the fast experiments ran in a shorter period than the slow experiments. From the graph it appears that the two plots are diverging exponentially, while both are still increasing with a positive concavity. This shows that neither are of $O(N)$ order, and it is additionally unlikely that either are at a $O(\log N)$ order.